Analyzing Genre-specific Thematic Patterns in Video Game Reviews

A Machine Learning Approach Using TF-IDF Vectors and Steam User-generated Tags

Nico Benz

3583917

nico.benz@studserv.uni-leipzig.de

Project report for the module

10-207-0101

Current Trends in Digital Humanities:

Computational Game Studies

/nicobenz/GameStudies-SteamPredictions

Computational Humanities

Institute of Computer Science

Leipzig University

September 11, 2023

Abstract

This paper focuses on computational game studies as an increasingly popular branch of digital humanities. Video game reviews, together with user-generated genre tags for each game, have been extracted from the video game platform Steam, resulting in an exhaustive corpus of more than one billion token, which provides a comprehensive basis for analysis. Machine learning models including Multinomial Naive Bayes, Logistic Regression and Random Forest, were trained on review texts along with the genre tags associated with the game that each particular review is belonging to. The findings show that these models are capable of predicting a games genre given review texts, highlighting the presence of underlying genre-specific thematic features embedded within these texts. Analysis of TF-IDF vector scores also show different ranking of tokens related to topics or game content within each genre, further supporting the hypothesis of genre-specific thematic features in video game reviews.

Keywords: Computational Game Studies, Natural Language Processing (NLP), Machine Learning (ML), Computational Linguistics (CL), Steam Reviews, TF-IDF, Classifier Model

Contents

List of Tables

List of Figures

1   Introduction

Game studies are a highly interdisciplinary field attracting researcher with various backgrounds including anthropology, sociology, psychology, philology and others (Aarseth, 2001). This diverse combination of involved disciplines is the first indicator that video games are not easily described or categorised. Video games offer a vast amount of aspects to study including their visuals, sounds, texts, mechanics, narrative strategies, perspectives and many more. Because of this, researchers in the past struggled to put video games into their field of study. They have tried analysing video games as films like a variant of cinema or as texts like a variant of literature. But doing so does not capture their essence enough to be able to fully analyse them (Aarseth, 2001). Aarseth, 1997 considers games cybertexts and ergodic literature, where the reader is considered a player and has to do non-trivial things to advance in the text and reach further parts of it (Aarseth, 1997, pp. 1, 4). In cybertexts access to parts of the text is based on the players decisions, actions or paths taken what results in a higher degree of narrative control that is not the case for traditional cinema or literature (Aarseth, 1997, pp. 3–4). Another important aspect in which games differ from films or literature is their inherent dependency on interactivity which makes them a non-linear and unstable kind of media (Apperley, 2006, pp. 6–7). Because of the mentioned reasons, games should be considered its own form of media with game studies as its own academic field (Aarseth, 2001).

Games differing in all these and many more aspects leads to the creation of genres and sub-genres that have several things in common while still being unique in their own ways, creating a vast field of game genres to analyse in game studies. The research of game genres was mostly dictated by research of genres in film and literature but because of the aforementioned challenges of treating games as movies or literature, the mapping of movie and literature genres on games does not take into account the unique properties of video games especially leaving out the important aspect of interactivity that should always play a role in video game genres (Apperley, 2006, p. 7).

This interactivity is probably the main connection between the player and the game. This paper assumes that player interactions with games of differing genres shows a statistical pattern of thematic properties that is dictated by the games genre. These genre-specific patterns can then be detected and measured in user-generated reviews as thematic pattern. The goal of this paper is to find evidence for these pattern by training machine learning classifiers on game reviews written by players and by analysing TF-IDF scores of the relevant tokens for each genre. Game reviews can be a way of accessing how players in general think about certain games Zagal et al., 2011, p. 358 and the results of this study could therefore be helpful to further investigate the culture of gaming and what players find most interesting in different genres.

2   Related work

2.1   Genres in video games

As already mentioned, in the past genres in video games have suffered from the fact that researchers tried to force genres of films or literature on them with a focus too strong on aesthetics instead on their ergodic nature (Apperley, 2006, p. 7). Heintz and Law, 2015 argue that even in the recent literature, video game genres lack clarity and consistency what can be seen in their definition on unrelated aspects like camera position as in first- or third-person shooter or intention as in educational games (Heintz & Law, 2015, pp. 176–177). Combined with the fact that games genres are not easy to distinct from

one another or from their sub-genres because of their fluidity, this makes games hard to categorise in genre labels (Heintz & Law, 2015, p. 177; Atkins, 2003, p. 23).

On Steam, the notion of genre is implemented differently. Users have the possibility to add new tags to games or upvote existing tags. This creates a list of attributes for each game on Steam that achieves a similar thing like traditional genres try to do: give information on the properties of a game. Simonson et al., 2023 investigated this approach and clustered games based on traditional genre labels given by game developers and compared them with user-generated tags from Steam. They could find that the traditional genre approach performed poor and clustered games together that did not share many features or properties and that user-generated tags did a better job in clustering related games together. However, they also found that the results using the user-generated tags were not nearly perfect. Windleharth et al., 2016 could also find evidence to support these findings and argue that user generated tags are in general more able to finely categorise media compared to curated genres or tags (Windleharth et al., 2016, p. 421).

2.2   Use of classifiers

Label classification tasks are thoroughly researched up to this date, including research questions similar to the one of this paper.

Bahuleyan, 2018 used labeled music tracks of the genres Hip Hop, Pop, Vocal, Rhythm, Reggae, Rock and Techno to train classifier models for a multi class destinction task in order to predict the genre of these songs. They could show that training classifiers like Logistic Regression, Random Forest and Support Vector Machines on feature lists of these tracks yields good prediction results. They also trained Convolutional Neural Networks and achieved F1 scores of 0.61 on spectrograms of the same labeled music tracks.

Kumar et al., 2022 trained classifier models on movie plots along with genre information and were able to achive very high F1 scores of above 0.9.

Gupta et al., 2019 trained classifier models on text data from books using TF-IDF vectors with similar results.

Göring et al., 2020 used gaming videos from games of different genres to train classifier models like Random Forest, Support Vector Machines, k-nearest neighbours or Gradient Boosting. They included the genres First-Person Shooter, Jump 'n' Run, RPG, Real Time Strategy, Top-Down Roleplay and Third-Person Shooter. The authors could show that the models were able to perform with F1 scores of up to 0.6.

Jiang and Zheng, 2023 used image-based, text-based and multimodal models trained on game covers and game descriptions various genres. They could show that model metrics were inconsistent throughout genres, ranging from 0.02 and 0.59 accuracy for image-based, 0.02 to 0.72 for text based and 0.08 to 0.72 for multimodal models.

These are just some examples of the usage of classifier models in order to predict the genre of media objects. All studies vary greatly in terms of their achieved F1 scores and other metrics, but it is also hard to compare their findings because all studies used a different amount of genres in their multi class classification task. This results in a differing amount of ways where the classifier could fail or perform suboptimally. When looking at the genres that most of these studies chose, it is interesting to see that most of these genres were very destinct from each other, no matter if game, music or movie genre.

All of these studies have in common though, that they chose different kinds of data to train their models, while non of them chose to use text reviews.

3   Methodology

As already mentioned, researchers in the past could show that machine learning models are very capable of performing class or genre classification in various situations and for various types of media. Because of that it can be expected that these classifiers will also perform well when trained on review data from games of different genres, given the assumption that there are genre-specific features like lexical choice or selection of thematic focus.

To get the data, reviews were collected from the Steam platform using their API as part of a custom script to boost collection rates. Simultaneously, their user-generated tags were collected from each games' main page using a custom crawler. Reviews were then randomly selected to create the training dataset. To avoid bias and possible confounds, only games with at least 10 reviews were part of the random selection and only reviews with a token count between 20 and 1000 after removal of stopwords and other steps of cleaning were regarded as viable and could end up in the dataset. Data collection finished after 50,000 reviews for every genre were collected.

In order to reduce noise from other confounding features like syntax or semantics, Term Frequency–Inverse Document Frequency (TF-IDF) vectors were chosen as training data for the classifier models. TF-IDF vectors are inherently better at showing lexical importance of terms in a collection of documents than they are capable of showing semantic relations within texts. For the goal of this paper this is an advantage in two ways: It reduces the noise that semantic patterns could create, and it also boosts the importance of lexical items and the thematic patterns that can be infered from their presence through TF-IDF scores. The three models Multinomial Naive Bayes, Logistic Regression and Random Forest were selected as models for the classification task. These models could show in the past that they are very capable of multi class classification tasks while not being overly computational intensive like neural networks or transformer models would be.

For verdicts on the papers assumptions both the models classification results and the TF-IDF scores will be analysed and evaluated.

4   Data overview

On the Steam platform users have the ability to write a public text review for any game. These reviews can be retrieved using the official Steam API. Given the vast amount of review data available on steam, this process took some time with a Python script generating API requests and saving the retrieved reviews non-stop for more than two months between May and July 2023. The user generated genre tags were not part of the data retrieved from the API and have been retrieved using a custom script crawling each Steam games web page and scraping the tags from the HTML content. Both processes were running in parallel but the crawling of the tags finished much faster, because of the very low amount of target data. This resulted in very small discrepancies between the review and the tag dataset because of games being removed from or added to Steam in the time window when the tag crawling was already finished but the review extraction not yet. Affected games where either the review set or the tag set was missing were not used in the combined corpus.

4.1 Games

In total, data of 134,076 games was downloaded which includes all games present on the platform at that time. Of these games all of their reviews have been downloaded. In some rare cases retrieval of all reviews of a certain game was not possible due to the API not giving a new cursor to fetch the next batch of reviews. After 10 tries of not getting a new cursor the script moved on to the next game, saving only the reviews gathered up to that point. This problem only occurred for games with several million reviews and only after at least one million reviews have been downloaded of that game. Because such great amounts of reviews for a single game would not be used for training to avoid strong bias or imbalance of the dataset anyway, this will not pose a challenge. See figure 1 for an overview of the distribution of games and review numbers.

Figure 1: Distribution of Reviews Across Games

4.2 Reviews

Among all these games the numbers of reviews for each game are not evenly distributed, as was expected. The biggest portion of games do not have reviews at all, while some games have more than one million reviews. These reviews also vary strongly in their length but all within the dimensions set by Steam which needs a review to be between 5 and 8,000 characters. The Steam API returns reviews as batches in a JSON format containing the reviews along with some anonymous metadata about the author and the review. For some basic metrics of the reviews see table 1.

	Total	Only English
Reviews	100,855,789	46,616,033
Tokens	15,023,225,453	10,643,950,482
Mean Review Size	148.96	228.33
Median Review Size	20	59

Table 1: Review Metrics

4.3 Genres

The genre term is used very loosely in this paper. Through the mechanic of user-generated tags, games have a list of user given labels along with the number of users that have given or upvoted a tag. If a game has more than 20 tags associated with it, only the 20 most common tags will be displayed on Steam. While most tags correspond to traditional genre labels (e.g. adventure, strategy, RPG) some of them focus more on features that could be present in any genre (e.g. early access, free to play, VR). There are also tags where it is not entirely clear whether they correspond to a genre or not (e.g. indie, exploration, casual) because of the fluidity of emerging new genres and subgenres. To avoid biased selection this paper will consider all user generated tags genres. Further sampling and clustering of these tags into more traditional genres is always possible for future research at a later date. See figure 2 for an overview of genres that are found within the 5 most common genres of all games in the corpus.

Figure 2: Number of Games for the Ten Most Common Genres

5 Experimental design

5.1 Corpus generation

The review data was gathered utilising Python scripts[1] along with the Steam API. The genre data was not part of the Steam API and was therefore crawled through a custom script accessing each games

---

[1]See the repository linked on the front page for access to all scripts and code used.

page on the Steam website and downloading its HTML content for scraping in a later step. Review and genre data has been combined by saving it in JSON files storing information like game id, review id, review text and most common genre tags for every review.

5.2 Dataset sampling

Only games were included that had at least 10 reviews and each review was randomly chosen and had to be between 20 and 1000 token while collecting maximum of 1000 reviews per game. Collected reviews where tokenised, cleaned from stopwords and custom stopwords, special characters and any other noise like ASCII art, and then lemmatised if possible. Custom stopwords were collected through comparison of most prominent tokens in earlier corpora using TF-IDF scores among all genres present in the corpora. The 50 most common tokens for all genres have been collected and counted. Among these tokens, only those present in at least three genres have been removed. Those consisted mostly of tokens belonging to the text genre review like *recommend*, *feel* or *like* or of very generic adjectives like *good*, *bad* or *fun*. Tokens belonging to the broad field of video games like *game* or *play* were also added to the custom stopwords. Content-specific words like *level* or *gameplay*, however, have not been removed. For a full list of all custom stopwords that were removed see appendix A. TF-IDF vectors were created off the cleaned reviews and then passed on to training. During early stages of dataset sampling the usage of Synthetic Minority Oversampling Technique (SMOTE) was planed in order to reduce small imbalances in the dataset that resulted from small differences in review lengths in each genre. Unfortunately SMOTE proved to be too computationally intensive to be implemented for a paper of this scope. Because of the general attention for a balanced dataset this will be of no major consequences.

5.3 Model training

The already prepared dataset was separated into TF-IDF vectors and genre labels, and an 80/20 split has been performed for separating the training from the test set. The models were trained using cross-validation in the form of k-folds with five folds. Various configurations in hyperparameter tuning were tested, but the standard values of the models implementations in Scikit-learn were nearly perfect. Random Forest was left on the standard values while setting Multinomial Naive Bayes alpha value to 0.5 and Logistic Regressions C value to 0.5 achieved best results. Metrics like precision, recall and F1 score for each fold were taken both in the form of mean values and genre-specific and will be reported within the next section. These metrics were also used to draw receiver operating characteristic (ROC) curves in a One-versus-Rest (OvR) fashion and precision recall curves for better model evaluation. After training, relevant metrics like accuracy, recall, precision and F1 score were collected.

6 Results and discussion

The models all performed relatively similar, with Random Forest being the least effective classifier. See table 2 for a result overview.

The relatively high values in F1 scores show that the models can predict the genre of the game from a review text way above baseline level of random guessing what would be at values of around 0.2. Another indicator of good model performance is the fact that precision, recall and F1 score are all very closely together, hinting at very stable model performance. Additionally, performance across

	Mean	Naive Bayes	Logistic Regression	Random Forest
Recall	0.66	0.68	0.68	0.62
Precision	0.66	0.68	0.68	0.62
F1 Score	0.66	0.68	0.68	0.62
Support	50000.0	50000.0	50000.0	50000.0

Table 2: Classifier Model Metrics Across All Genres

all five folds is very similar what further enhances the claim of stable and well performing models. See table 3 for the metrics for every fold.

	Average	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5
Naive Bayes						
Recall	0.68	0.55	0.7	0.69	0.74	0.74
Precision	0.68	0.62	0.71	0.66	0.66	0.75
F1 Score	0.68	0.58	0.71	0.67	0.7	0.74
Support	50000.0	9912.0	10200.0	9914.0	10036.0	9938.0
Logistic Regression						
Recall	0.68	0.57	0.71	0.68	0.7	0.75
Precision	0.68	0.59	0.71	0.67	0.71	0.72
F1 Score	0.68	0.58	0.71	0.67	0.7	0.74
Support	50000.0	9912.0	10200.0	9914.0	10036.0	9938.0
Random Forest						
Recall	0.62	0.51	0.62	0.63	0.64	0.69
Precision	0.62	0.5	0.66	0.6	0.65	0.69
F1 Score	0.62	0.5	0.64	0.61	0.64	0.69
Support	50000.0	9912.0	10200.0	9914.0	10036.0	9938.0

Table 3: Classifier Model Metrics Across All Folds

When delving further into analysis of model performance, one should always look at performance differences between the predicted labels. Table 4 shows the models performance across all tested game genres along with a mean value across all genres.

Interestingly, genres do not perform on the same level. While most genres perform on a similar level, adventure games are harder to predict for the models than puzzle games, having a difference in F1 score of 0.16 in Multinomial Naive Bayes and with similar values in the other two models. I argue that this lies in the nature of the adventure genre that shows more overlap with other genres and more individual fluidity for every instance of adventure games. These games contain elements of other genres to a differing degree and also have fuzzy boundaries to the genres like action-adventure or RPG. This can be further observed in the ROC and precision recall curves in figure 3.

The adventure genre has the lowest area under curve (AUC) from all genres in their respective ROC curves, even across all models, indicating that this is not because of model performance issues but a problem inherent to the genre. The precision recall curves show the lowest average precision (AP) for adventure and also a stronger initial decline in performance for this genre alone, also across all models with the steepest fall in the Random Forest classifier. The curve shows the trade-off relation of how many predictions are actually of the predicted label (precision) and how many instances of a label are correctly predicted (recall). The strong initial decline compared to other genres shows that the model struggles especially with positive predictions of the adventure genre what I regard as evidence for either a stronger distraction from other genre labels or missing predictive power in the adventure class.

9

	Average	Adventure	Strategy	Simulation	RPG	Puzzle
Naive Bayes						
Recall	0.68	0.55	0.69	0.69	0.73	0.74
Precision	0.68	0.6	0.71	0.66	0.66	0.74
F1 Score	0.68	0.58	0.7	0.67	0.69	0.74
Support	50000.0	10000.0	10000.0	10000.0	10000.0	10000.0
Logistic Regression						
Recall	0.68	0.57	0.7	0.68	0.69	0.75
Precision	0.68	0.58	0.7	0.66	0.71	0.72
F1 Score	0.68	0.57	0.7	0.67	0.7	0.74
Support	50000.0	10000.0	10000.0	10000.0	10000.0	10000.0
Random Forest						
Recall	0.62	0.51	0.61	0.63	0.63	0.69
Precision	0.62	0.49	0.66	0.59	0.65	0.69
F1 Score	0.62	0.5	0.63	0.61	0.64	0.69
Support	50000.0	10000.0	10000.0	10000.0	10000.0	10000.0

Table 4: Classifier Model Metrics Across All Genres

To further investigate the nature of the failing prediction within this class, one should pay attention to the TF-IDF scores for every genre to see the most prominent tokens for every genre. See table 5 for an overview.

Adventure	Strategy	Simulation	RPG	Puzzle	
story, 0.25	units, 0.15	route, 0.15	story, 0.2	levels, 0.22	
short, 0.13	strategy, 0.14	vr, 0.15	combat, 0.16	level, 0.17	
music, 0.11	tower, 0.12	work, 0.12	characters, 0.13	story, 0.16	
gameplay, 0.11	ai, 0.11	better, 0.11	character, 0.12	music, 0.14	
characters, 0.11	gameplay, 0.11	buy, 0.11	hours, 0.11	simple, 0.12	
graphics, 0.1	cards, 0.1	real, 0.1	gameplay, 0.11	easy, 0.12	
level, 0.1	campaign, 0.1	things, 0.09	system, 0.1	gameplay, 0.11	
experience, 0.1	defense, 0.1	money, 0.09	better, 0.1	short, 0.11	
art, 0.1	turn, 0.1	price, 0.09	far, 0.09	hours, 0.1	
point, 0.09	hours, 0.1	experience, 0.09	world, 0.09	hard, 0.09	
interesting, 0.09	card, 0.1	free, 0.08	buy, 0.09	art, 0.09	
controls, 0.09	better, 0.09	graphics, 0.08	interesting, 0.09	achievements, 0.09	
price, 0.09	buy, 0.09	train, 0.08	content, 0.08	price, 0.09	
better, 0.08	player, 0.09	add, 0.08	enemies, 0.08	mechanics, 0.09	
character, 0.08	war, 0.08	sim, 0.08	things, 0.08	challenging, 0.09	
⋮	⋮	⋮	⋮	⋮	⋮

Table 5: Tokens with TF-IDF Scores by Genre

The table shows the most relevant tokens for every genre along with their TF-IDF score. While most genres have one token standing out with a value a bit higher than other tokens, the highest scoring token for adventure is scored nearly twice as high as the following token. That indicates that the token *story* is the main predictor for the genre adventure. Having a good predictor is usually a good sign for a class but the problem for the adventure genre arises especially in the fact that there are no other strong predictors and that their main predictor is not unique for this genre while the other genres have strong and especially unique predictors. This will get more clear when looking at the ten highest mean TF-IDF scores in table 6.

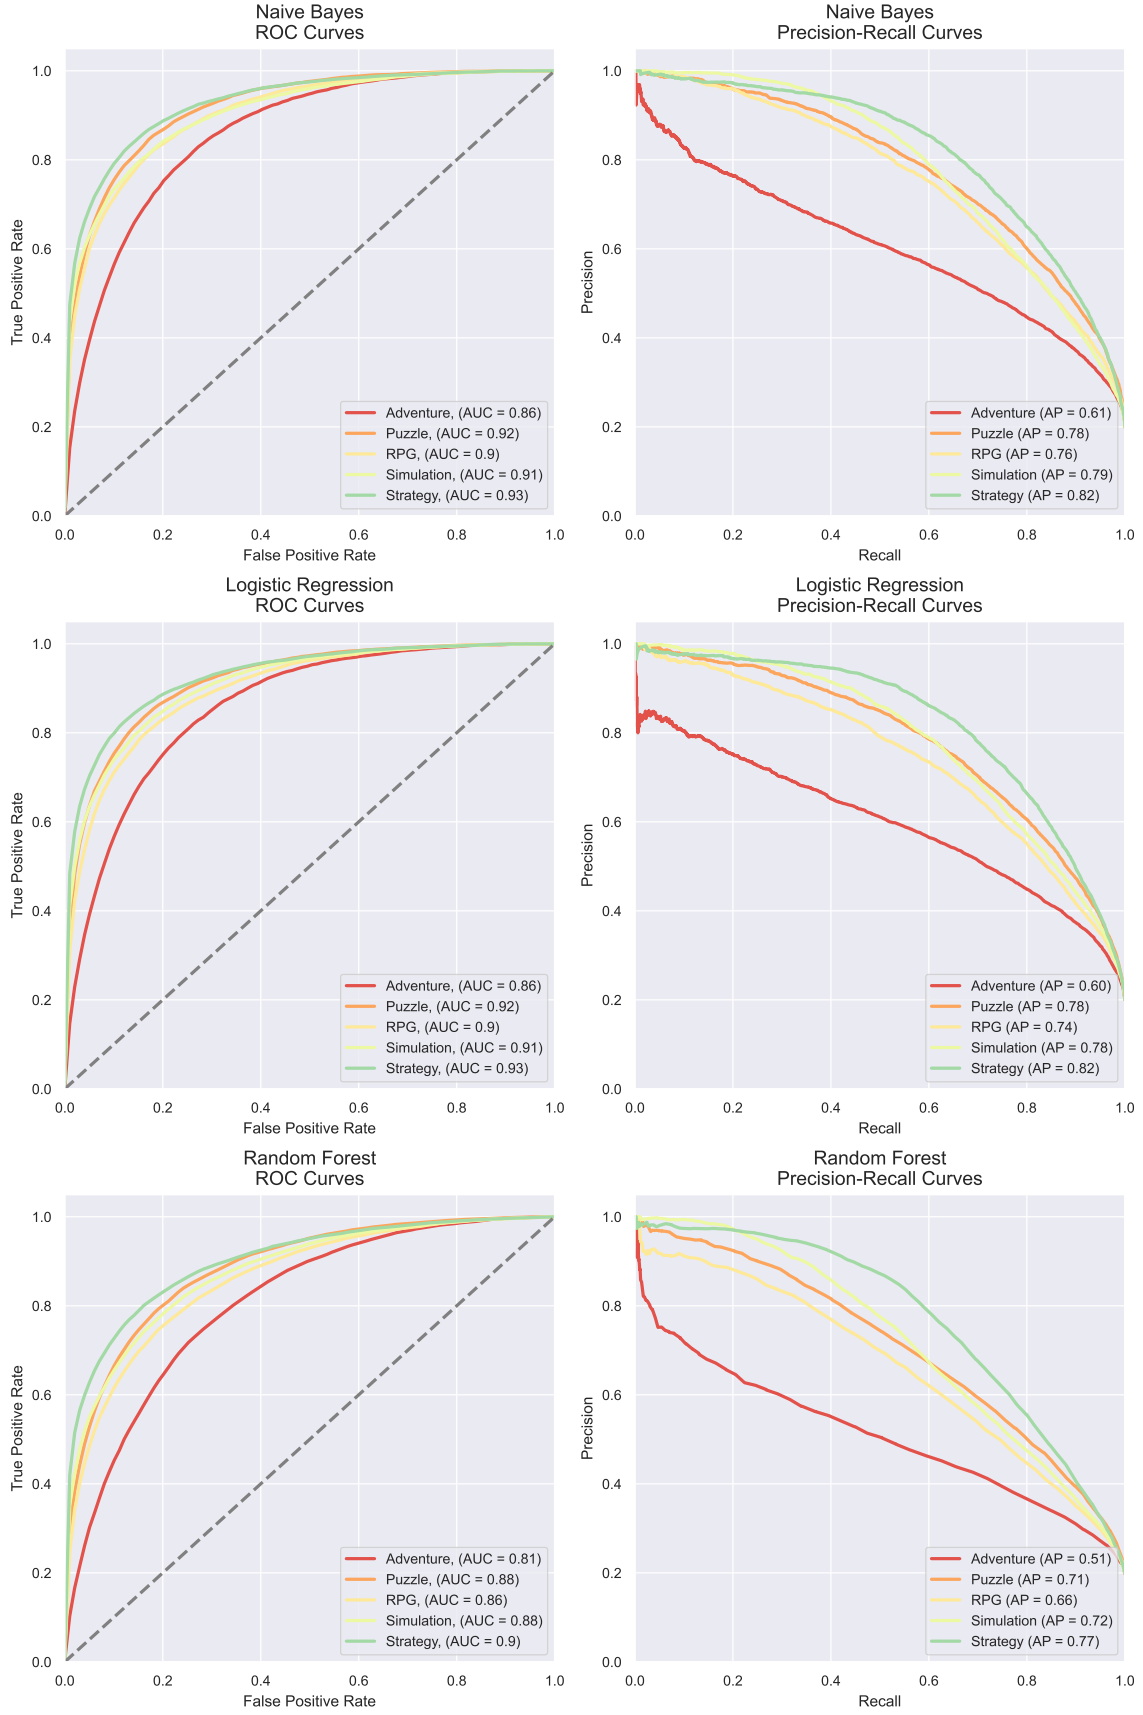It is clearly visible that *story* is the most important token across all genres when looking at the

Figure 3: ROC OvR Curves and Precision Recall Curves for All Models

mean TF-IDF scores across all classes since it is present in all classes but strategy. It could be easily assumed that classes that also have high TF-IDF scores for the token *story* like RPG and puzzle are

Token	Adventure	Strategy	Simulation	RPG	Puzzle
story	0.25	-	0.06	0.2	0.16
gameplay	0.11	0.11	0.06	0.11	0.11
hours	0.08	0.1	0.07	0.11	0.1
better	0.08	0.09	0.11	0.1	0.07
level	0.1	0.08	-	0.08	0.17
buy	0.07	0.09	0.11	0.09	0.06
graphics	0.1	0.08	0.08	0.07	0.08
price	0.09	0.07	0.09	0.07	0.09
work	0.07	0.08	0.12	0.07	0.06
things	0.08	0.07	0.09	0.08	0.07

Table 6: Tokens with TF-IDF Scores Across Genres Sorted by Mean Score

the main distractors for this class and cause the most false positives. Interestingly, looking at the confusion matrix in figure 4 reveals that actually the opposite is the case.


Figure 4: Distribution of Reviews Across Games

The highest amount of false positives for the adventure genre is caused by the strategy genre, followed by the simulation genre. The other two genres that also have a high TF-IDF score with the token *story* are responsible for the least amount of false positives. While this seems counter-intuitive at first glance, I argue that the answer for that is quite easy. The false positives are most probably causes by review texts that do not deal with the topic of story and avoid the token. Without its main predictor, a review text for a game of the adventure genre could easily be regarded as a review for a game of the strategy or simulation genre, mostly because reviews in these genres are characterised through the statistical lack of mentioning the story of a game. When looking at the other genre that

has high TF-IDF scores for the token *story*, RPG, the problem does not occur since the confusion matrix indicates that strategy and simulation are responsible for the least amount of false positives for the genre RPG. This can be explained by the fact that this genre, while having *story* as their highest ranked TF-IDF score, also has strong other predictors that also are not ranked high for the other genres. The distribution of RPGs predictive power is not as strongly centered on one token like in the case of adventure and RPG has other strong predictors like *combat* that are way less relevant for the other genres analysed. This creates further evidence for my claim that the low performance of the adventure genre is caused by a lack of unique predictors that other genres have like *units* for strategy, *route* for simulation, *combat* for RPG or *level* for puzzle. I regard the existance of these differing TF-IDF scores to be genre-specific thematic patterns. The data shows that some tokens like *graphics* or *price* show a rather similar distribution across genres, while other token like *story* show a higher score in adventure and RPG genre. There is also evidence for topics that are relevant to most genres but less relevant to one genre, like in the case of *gameplay* that has a very similar score across all genres except simulation, where it is scored much lower.

However, there are still some limitations of this study that are needed to be mentioned. First of all, the dataset with its size of 250,000 review texts is rather small compared to the size of the whole corpus of several million review texts. Future research should be done including more review texts of the corpus. In terms of genres, only five genres where used in this study while there are several hundred in existence. Even if only the most relevant genres were considered, this would still end up with about four times as much genres as used in this study. When using more genres, future research should also consider not excluding games with overlapping genres or make it a multi label classification task in the first place where models are trained on more than one genre label for each review text. In this context the use of weighted genre labels, generated through the number of user that upvoted a certain genre tag, could be used for less distortion caused by including many labels.

7   Conclusion

For this paper I trained three different machine learning classifiers on review texts from the platform Steam and analysed TF-IDF vectors created from these review texts in order to investigate thematic differences between genres. I could show that models like Multinomial Naive Bayes, Logistic Regression and Random Forest are in fact capable of achieving reasonable results when predicting genres learned from user-generated reviews texts along with user-generated genre tags. I could further show that there are thematic patterns that are genre-specific by analysing TF-IDF vectors and their scores between genres. While there are some topics like *gameplay*, *graphics* and *price* that seem to have the same relevance to players across the tested genres, other themes like *story* seem to matter way more to players of adventure games than to players of strategy or simulation games. Future research has to see whether these findings can be replicated on other genres then the ones analysed in this study.

References

Aarseth, E. (1997). *Cybertext: Perspectives on ergodic literature*. JHU Press.

Aarseth, E. (2001). Computer game studies, year one. *Game studies*, *1*(1), 1–15.

Apperley, T. H. (2006). Genre and game studies: Toward a critical approach to video game genres. *Simulation & Gaming*, *37*(1), 6–23. https://doi.org/10.1177/1046878105282278

Atkins, B. (2003). *More than a game*. Manchester University Press. https://doi.org/10.7228/manchester/9780719063640.001.0001

Bahuleyan, H. (2018). Music genre classification using machine learning techniques. https://doi.org/10.48550/ARXIV.1804.01149

Göring, S., Steger, R., Rao Ramachandra Rao, R., & Raake, A. (2020). Automated genre classification for gaming videos. *2020 IEEE 22nd International Workshop on Multimedia Signal Processing (MMSP)*, 1–6. https://doi.org/10.1109/MMSP48831.2020.9287122

Gupta, S., Agarwal, M., & Jain, S. (2019). Automated genre classification of books using machine learning and natural language processing. *2019 9th International Conference on Cloud Computing, Data Science & Engineering (Confluence)*. https://doi.org/10.1109/confluence.2019.8776935

Heintz, S., & Law, E. L.-C. (2015). The game genre map: A revised game classification. *Proceedings of the 2015 Annual Symposium on Computer-Human Interaction in Play*, 175–184. https://doi.org/10.1145/2793107.2793123

Jiang, Y., & Zheng, L. (2023). Deep learning for video game genre classification. *Multimedia Tools and Applications*, *82*(14), 21085–21099. https://doi.org/10.1007/s11042-023-14560-5

Kumar, S., Kumar, N., Dev, A., & Naorem, S. (2022). Movie genre classification using binary relevance, label powerset, and machine learning classifiers. *Multimedia Tools and Applications*, *82*(1), 945–968. https://doi.org/10.1007/s11042-022-13211-5

Simonson, R. J., Keebler, J. R., & Doherty, S. (2023). The need for recategorized video game labels: A quantitative approach. *Game Studies*, *23*(1).

Windleharth, T. W., Jett, J., Schmalz, M., & Lee, J. H. (2016). Full steam ahead: A conceptual analysis of user-supplied tags on steam. *Cataloging & Classification Quarterly*, *54*(7), 418–441. https://doi.org/10.1080/01639374.2016.1190951

Zagal, J. P., Tomuro, N., & Shepitsen, A. (2011). Natural language processing in game studies research. *Simulation & Gaming*, *43*(3), 356–373. https://doi.org/10.1177/1046878111422560

Appendix A   Custom stopwords

Custom stopwords that were removed from the corpora before model training:

game, like, good, games, time, play, fun, way, great, little, bit, lot, pretty, feel, think, recommend, playing, things, want, different, played, worth, got, love, better, new, need, find, bad, nice, steam, know, dlc, use, hours, people, nt[2], adventure, strategy, simulation, rpg, puzzle, better, buy, thing, things

---

[2]This is probably an artifact created by stemming the falsely written *dont* (without apostrophe).